BLG435E Artificial Intelligence





Lecture 3: Problem-Solving





Outline



- Problem-solving agents
- Problem formulation
- Uninformed search strategies
- Informed search strategies
 - Best-first search
 - A* search
 - Heuristics



Problem Solving



- Goal formulation
 - limiting objectives
 - Uses current situation and agent's performance measure
 - A set of world states

- Problem formulation
 - What actions and states to consider

- Search procedure
 - Problem → solution in the form of action sequence



Problem-solving Agents



```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent: seq, an action sequence, initially empty
               state, some description of the current world state
               goal, a goal, initially null
              problem, a problem formulation
  state \leftarrow \text{UPDATE-STATE}(state, percept)
  if seq is empty then
      goal \leftarrow FORMULATE-GOAL(state)
      problem \leftarrow FORMULATE-PROBLEM(state, goal)
      seq \leftarrow SEARCH(problem)
      if seq = failure then return a null action
  action \leftarrow FIRST(seq)
  seq \leftarrow REST(seq)
  return action
```



Well-defined Problems



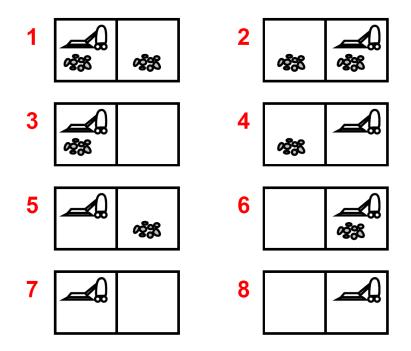
- Initial state
- Actions
- Transition model
- State space
 - Path in the state space
 - Size of the search space
- Goal test
 - Explicit vs. implicit
- Path cost
 - Step cost
- The solution to the problem
 - Optimal solution
- Abstraction



Vacuum-world Problem



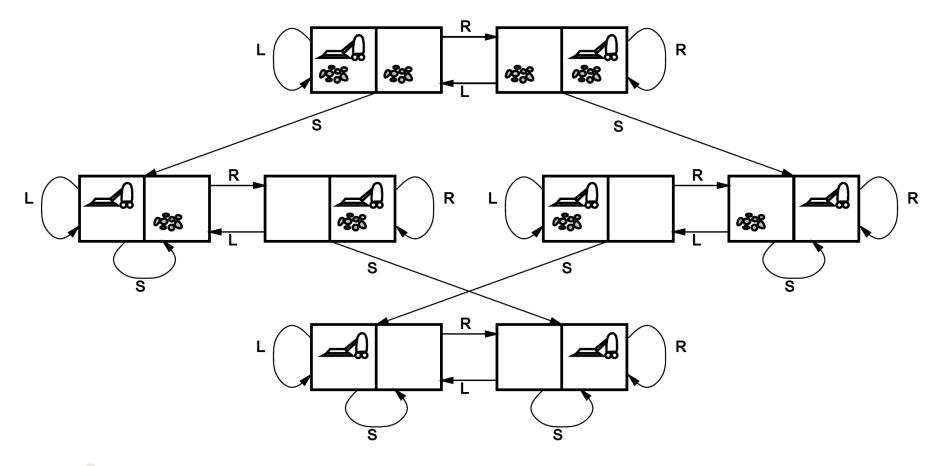
- Single state, start in #5
- Solution?





Vacuum-world State Space Graph

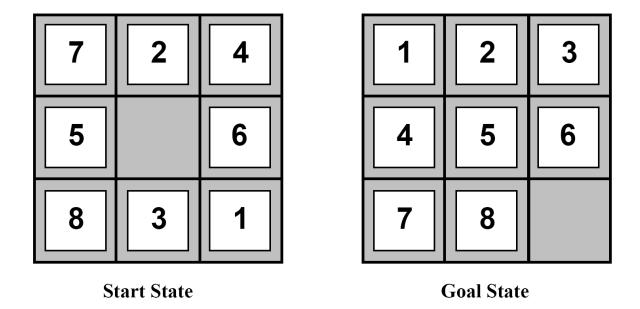






8-Puzzle

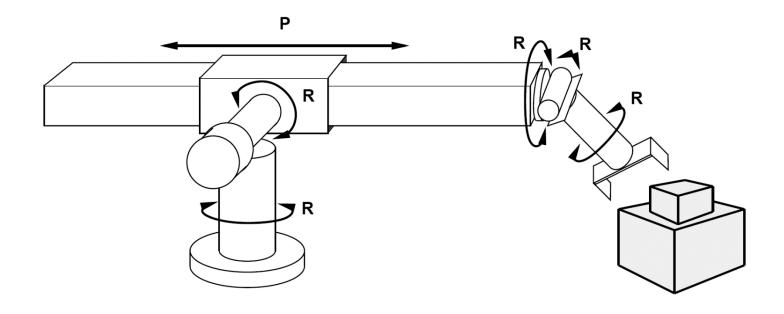






Robotic Assembly







Romania Problem

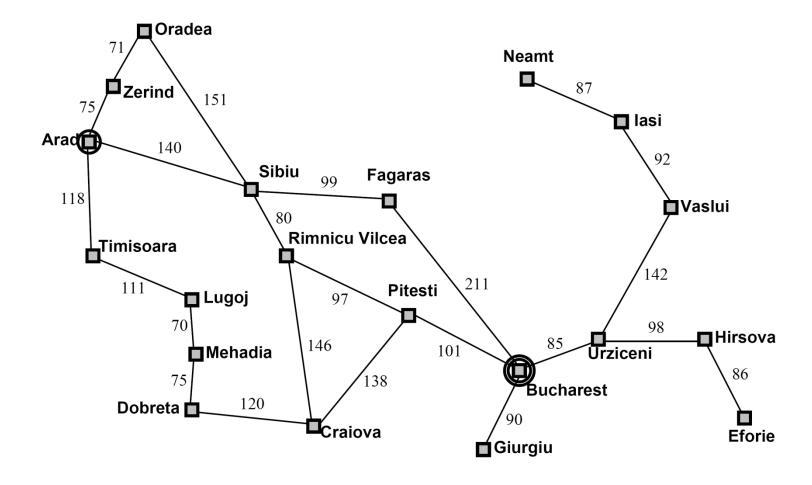


- On holiday in Romania, currently in Arad. Flight leaves tomorrow from Bucharest
- Formulate goal:
 - Be in Bucharest
- Formulate problem:
 - States: various cities
 - Actions: drive between cities
- Find solution: sequence of actions



Romania

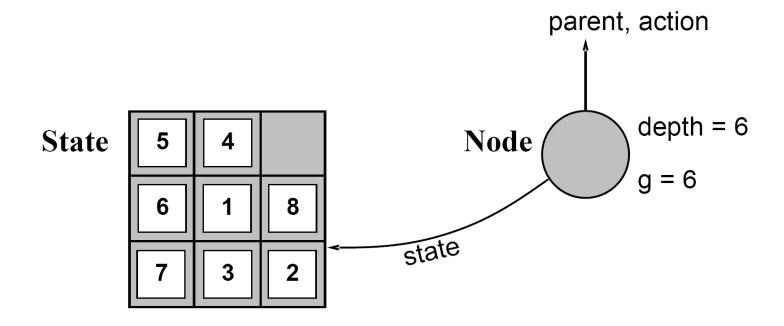






Implementation: states vs. nodes







Node data structure



- State: the state to which the node corresponds
- Parent-node: the node in the search tree that generated this node
- Action: action that was applied
- Path-cost: the cost from the initial state to the node g(n)
- Depth: the number of steps along the path from the initial state



Node data structure



Nodes = states ?

 Frontier: the collection of nodes that have been generated but not yet expanded

Implementation of frontier? Set?



Tree and Graph -Search Algorithms



Off-line, simulated exploration of the state space by expanding states

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
   if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
   if the node contains a goal state then return the corresponding solution
   expand the chosen node, adding the resulting nodes to the frontier
```

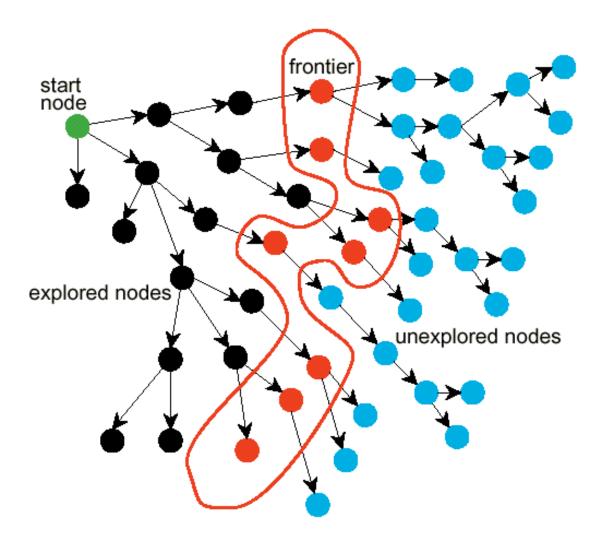
function GRAPH-SEARCH(problem) returns a solution, or failure
 initialize the frontier using the initial state of problem
 initialize the explored set to be empty
loop do
 if the frontier is empty then return failure
 choose a leaf node and remove it from the frontier
 if the node contains a goal state then return the corresponding solution
 add the node to the explored set
 expand the chosen node, adding the resulting nodes to the frontier
 only if not in the frontier or explored set

Choosing, testing, expanding



Nodes in The Search Tree







Measuring Problem-Solving Performance



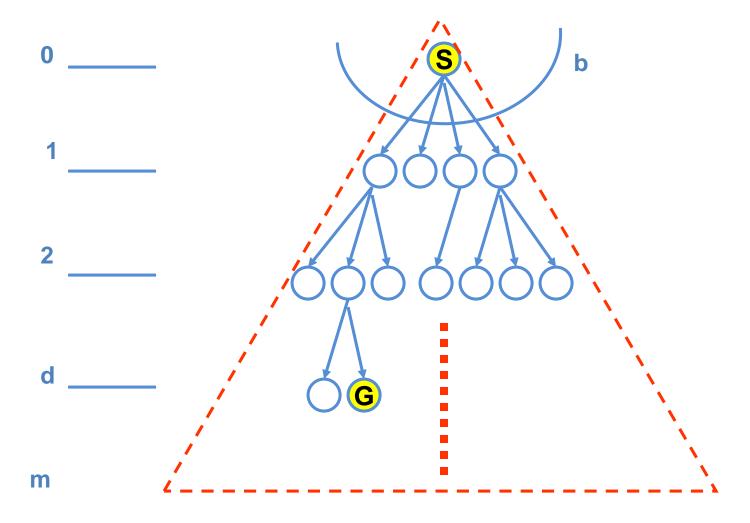
Strategy = order of node expansion

- Completeness
- Optimality
- Time complexity
- Space complexity
 - b: maximum branching factor
 - d: depth of the least-cost solution
 - m: maximum depth of the search tree



Measuring Problem-Solving Performance







BLG435E Artificial Intelligence





Uninformed Search Strategies





Uninformed Search Strategies



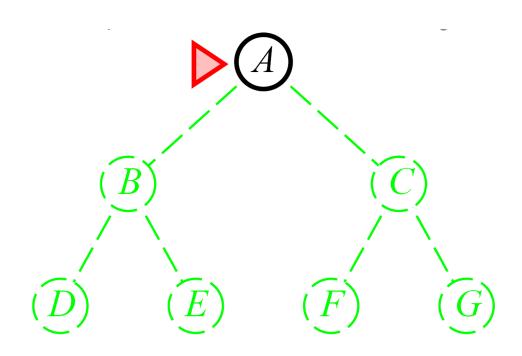
Use information only available in the problem definition

- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative Deepening search





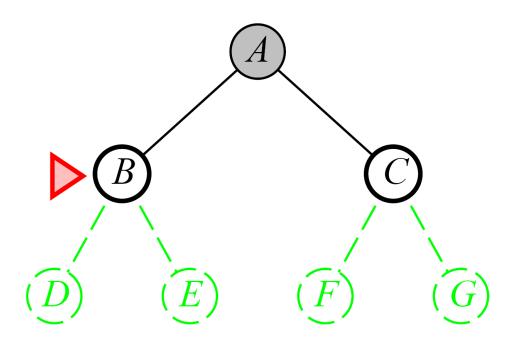
- Expand shallowest unexpanded node
- Frontier is a FIFO queue








- Expand shallowest unexpanded node
- Frontier is a FIFO queue








- Expand shallowest unexpanded node
- Frontier is a FIFO queue







- Expand shallowest unexpanded node
- Frontier is a FIFO queue





Properties of BFS



- Complete?
- Optimal?
- Time complexity?
- Space complexity?



Properties of BFS



- Complete? Yes, if b is limited
- Optimal? Optimal, if path cost is nondecreasing
- Time complexity? O(b^d)
- Space complexity? O(b^d) keeps every node in memory



BFS – Time/Space Requirements



• b = 10, 1.000.000 can be generated per sec

| Depth | Nodes | | Time | N | Memory | |
|-------|-----------|-----|--------------|------|-----------|--|
| 2 | 110 | .11 | milliseconds | 107 | kilobytes | |
| 4 | 11,110 | 11 | milliseconds | 10.6 | megabytes | |
| 6 | 10^{6} | 1.1 | seconds | 1 | gigabyte | |
| 8 | 10^{8} | 2 | minutes | 103 | gigabytes | |
| 10 | 10^{10} | 3 | hours | | terabytes | |
| 12 | 10^{12} | 13 | days | 1 | petabyte | |
| 14 | 10^{14} | | years | 99 | petabytes | |
| 16 | 10^{16} | | years | | exabytes | |

A node requires 1000 bytes of storage



Uniform-cost Search (UCS)



- Expand least-cost (g(n)) unexpanded node
- Frontier = queue ordered by path cost
- What if all costs are the same?

- Complete?
- Optimal?
- Time complexity?
- Space complexity?



Uniform-cost Search (UCS)




- Expand least-cost unexpanded node
- Frontier = queue ordered by path cost
- What if all costs are the same?

- Complete? Yes, if step cost > ϵ
- Optimal? Yes,
- Time complexity? O(b¹+LC*/ε」)
- Space complexity? O(b¹+LC*/ε






- Expand deepest unexpanded node
- Frontier = LIFO queue








- Expand deepest unexpanded node
- Frontier = LIFO queue








- Expand deepest unexpanded node
- Frontier = LIFO queue








- Expand deepest unexpanded node
- Frontier = LIFO queue







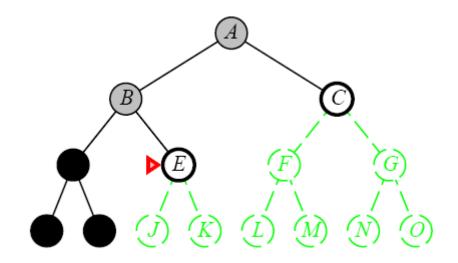
- Expand deepest unexpanded node
- Frontier = LIFO queue







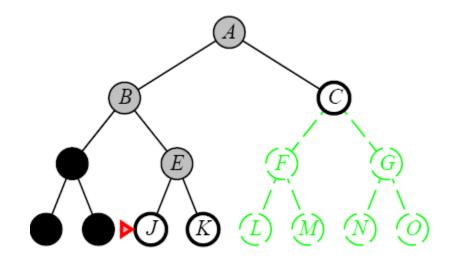
- Expand deepest unexpanded node
- Frontier = LIFO queue







- Expand deepest unexpanded node
- Frontier = LIFO queue

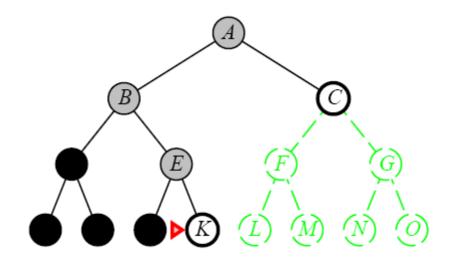




Depth-first Search (DFS)



- Expand deepest unexpanded node
- Frontier = LIFO queue

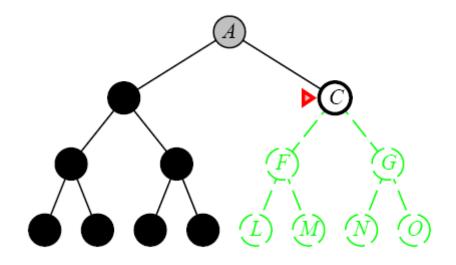




Depth-first Search (DFS)



- Expand deepest unexpanded node
- Frontier = LIFO queue

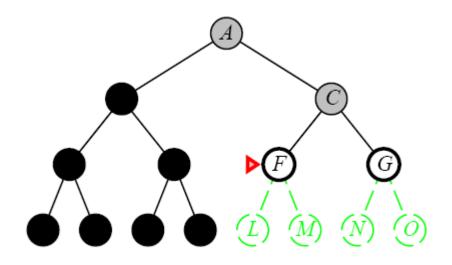




Depth-first Search



- Expand deepest unexpanded node
- Frontier = LIFO queue

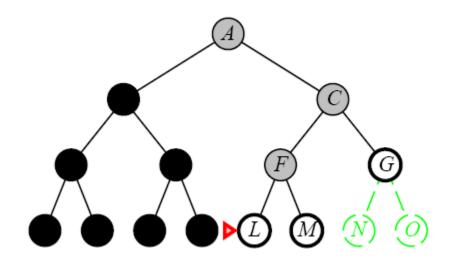




Depth-first Search



- Expand deepest unexpanded node
- Frontier = LIFO queue

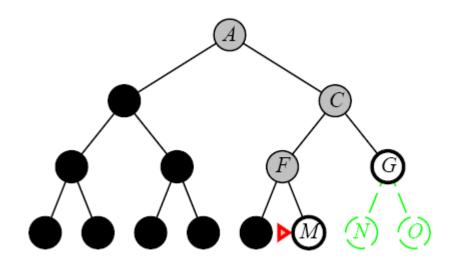




Depth-first Search



- Expand deepest unexpanded node
- Frontier = LIFO queue





Properties of DFS



- Complete?
- Optimal?
- Time complexity?
- Space complexity?



Properties of DFS



- Complete? No
- Optimal? No
- Time complexity? O(b^m)
- Space complexity? O(bm)



Depth-limited Search



Depth-first search with depth limit: l

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
cutoff-occurred? ← false
if GOAL-TEST(problem, STATE[node]) then return node
else if DEPTH[node] = limit then return cutoff
else for each successor in EXPAND(node, problem) do
result ← RECURSIVE-DLS(successor, problem, limit)
if result = cutoff then cutoff-occurred? ← true
else if result ≠ failure then return result
if cutoff-occurred? then return cutoff else return failure
```





```
function ITERATIVE-DEEPENING-SEARCH (problem) returns a solution inputs: problem, a problem for depth \leftarrow 0 to \infty do result \leftarrow \text{DEPTH-LIMITED-SEARCH}(problem, depth) if result \neq \text{cutoff then return } result end
```





Limit = 0




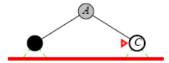









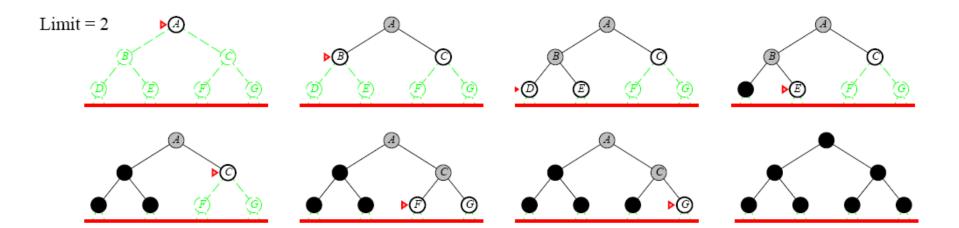






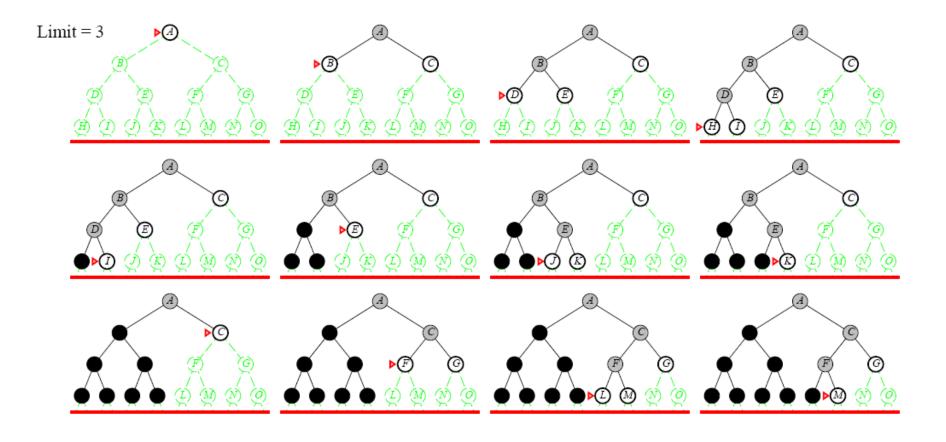














Properties of IDS



- Complete?
- Optimal?
- Time complexity?
- Space complexity?



Properties of IDS



- Complete? Yes, if b is limited
- Optimal? Yes, if path cost is nondecreasing
- Time complexity? O(b^d)
- Space complexity? O(bd)



Summary of Algorithms

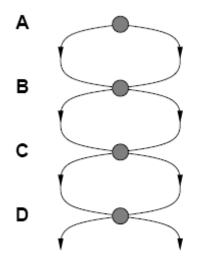


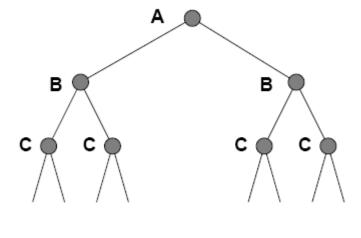
| Criterion | Breadth- | Uniform- | Depth- | Depth- | Iterative |
|-----------|----------|---------------------------------|--------|--------------------|-----------|
| | First | Cost | First | Limited | Deepening |
| Complete? | Yes* | Yes* | No | Yes, if $l \geq d$ | Yes |
| Time | b^d | $b^{\lceil C^*/\epsilon ceil}$ | b^m | b^l | b^d |
| Space | b^d | $b^{\lceil C^*/\epsilon ceil}$ | bm | bl | bd |
| Optimal? | Yes^* | Yes | No | No | Yes^* |



Avoiding Repeating States









General Graph-search Algorithm



function GRAPH-SEARCH(problem) returns a solution, or failure initialize the frontier using the initial state of problem initialize the explored set to be empty loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution add the node to the explored set expand the chosen node, adding the resulting nodes to the frontier only if not in the frontier or explored set



BLG435E Artificial Intelligence





Informed Search Strategies





Tree-Search Algorithms-Revisited



function TREE-SEARCH(problem) returns a solution, or failure initialize the frontier using the initial state of problem loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution expand the chosen node, adding the resulting nodes to the frontier

function GRAPH-SEARCH(problem) returns a solution, or failure
initialize the frontier using the initial state of problem
initialize the explored set to be empty
loop do

if the frontier is empty then return failure choose a leaf node and remove it from the frontier if the node contains a goal state then return the corresponding solution add the node to the explored set expand the chosen node, adding the resulting nodes to the frontier only if not in the frontier or explored set



Best-first Search



- Node is selected based on f(n)
 - Estimate of desirability
 - The node with lowest evaluation is selected

Frontier is a queue sorted in decreasing order of desirability

- A key component: heuristic function
 - -h(n) = estimated cost of the cheapest path from node n to a goal node



Greedy Best-First Search



Expands the node that is closest to the goal

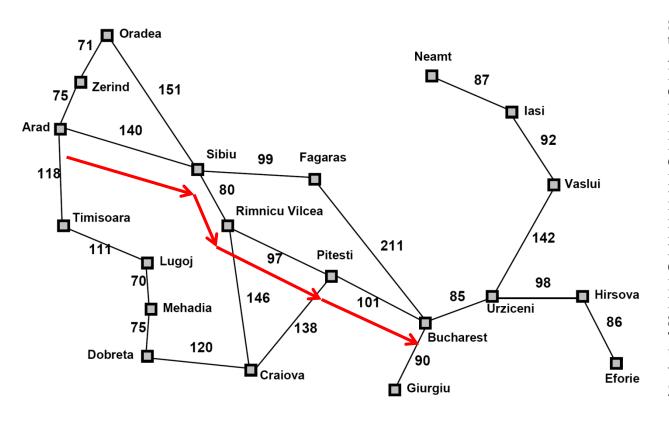
•
$$f(n) = h(n)$$

• $h_{SLD}(n) =$ Straight-line distance heuristic



Romania Problem Step Costs





| Straight–line distan to Bucharest | .ce |
|--------------------------------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |
| | 377 |































Romania Problem – Greedy Solution





| Straight-line distant | ice |
|-----------------------|-------|
| to Bucharest | |
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |
| | _ , . |



Properties of Greedy Search



- Complete? No, can stuck in loops
- Time? O(b^m), can be improved with a good heuristic
- Space? O(b^m)
- Optimality? No



A* Search



- The most well-known form of best-first search
- Avoiding paths that are already expensive
- f(n) = g(n) + h(n)- g(n) = the path cost from the start node to node n

• f(n) becomes the estimated cost of the cheapest solution through n



Admissible Heuristic

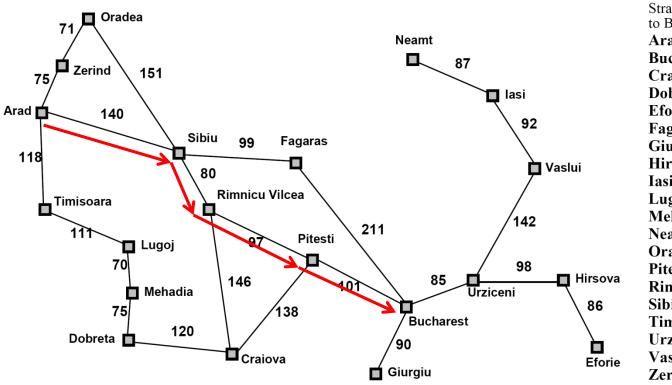


- Admissible heuristic:
 - $-h(n) \le h(n)^*$ (true cost from n)
 - $-h(n) \ge 0$ and h(G) = 0
- A* (with tree-search) is optimal if h(n) is admissible

• $h_{SLD}(n)$ admissible?



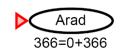




| Straight-line distant | ice |
|-----------------------|-----|
| to Bucharest | |
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 178 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |



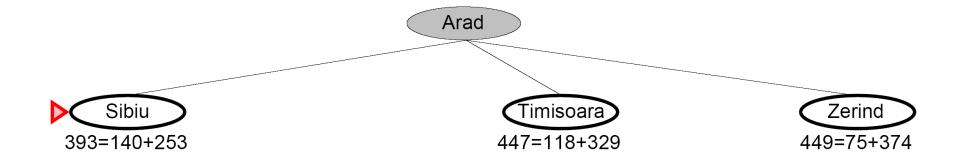




- g(n) values-computed from the step costs
- $h_{SLD}(n)$ values-known initially








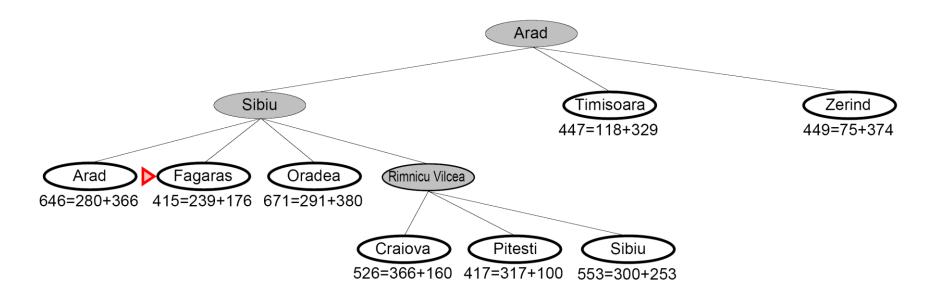






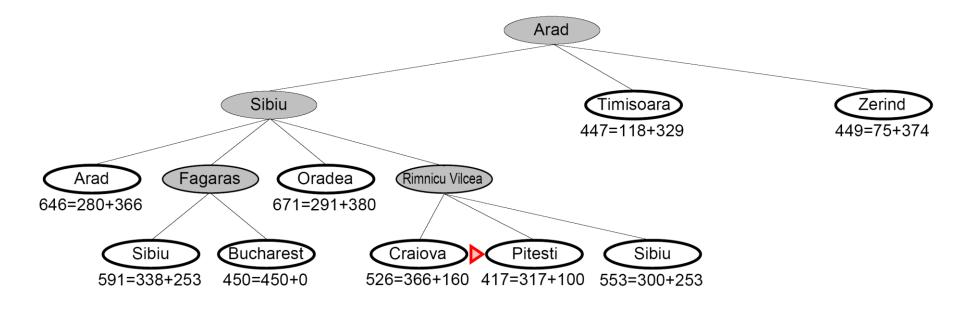






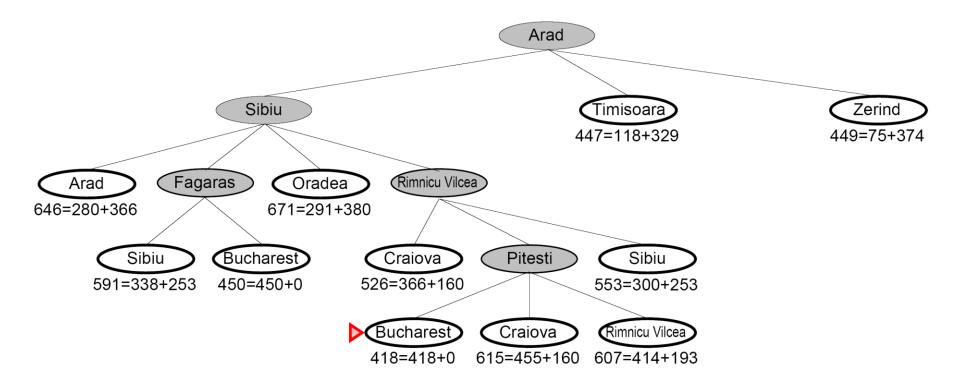










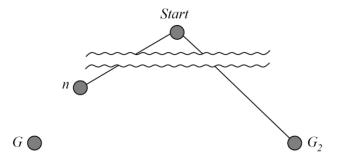




Optimality of A* (standard proof)



- Suppose a suboptimal goal G_2
- Let n be an unexpanded node on a shortest path to an optimal goal G_1



$$f(G_2) = g(G_2)$$
 since $h(G_2) = 0$
> $g(G_1)$ since G_2 is suboptimal
 $\geq f(n)$ since h is admissible

• Since $f(G_2) > f(n)$, A* will never select G_2 for expansion



Consistency



A heuristic is consistent if

$$h(n) \le c(n, a, n') + h(n')$$

If h is consistent, we have

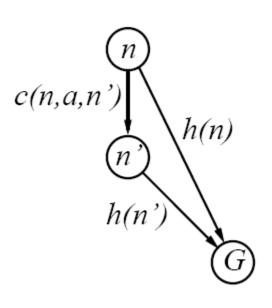
$$f(n') = g(n') + h(n')$$

$$= g(n) + c(n, a, n') + h(n')$$

$$\geq g(n) + h(n)$$

$$= f(n)$$

I.e., f(n) is nondecreasing along any path.





When do we need a consistent heuristic?



 We need an admissable and consistent heuristic for the graph search version of A*

Why?



Properties of A*



- Complete? Yes, unless there are infinitely many nodes with f < C*
- Time? Exponential in
 - [relative error in h * length of solution]
- Space? Keeps all nodes in memory
- Optimality? Yes



To reduce the space complexity



- Iterative deepening A* (IDA*)
 - Cut-off used is the f-cost

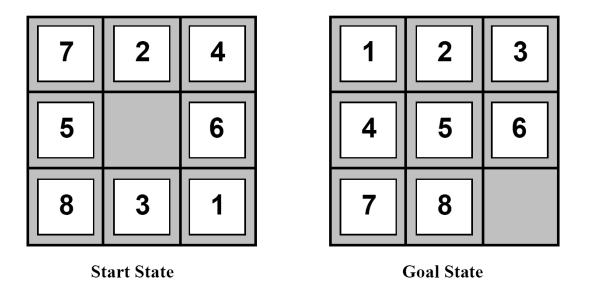
- Recursive best-first search (RBFS)
 - f_limit of best alternative path
- Memory-Bounded A* (MA*)
- Simplified MA* (SMA*)
 - drops the worst leaf node when memory is full



Admissible Heuristics



• 8-puzzle



- $h_1(n)$ = the number of misplaced tiles
- $h_2(n)$ = the total Manhattan distance



Dominance



• if $h_2(n) \ge h_1(n)$ for all n (both admissible) then h_2 dominates h_1 , and is better for search

• Given any admissible heuristics h_a, h_b

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible and dominates $h_a,\,h_b$

