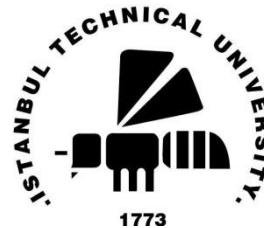


BLG435E

Artificial Intelligence



Lecture 4: Constraint Satisfaction Problems



- CSP problem formulation
- CSP examples
- Backtracking search for CSPs
- Problem structure and problem decomposition
- Local search for CSPs

Constraint Satisfaction Problems

- Search algorithms so far:
 - state is a “black box”
 - domain specific heuristics
 - states are accessible by problem specific routines
- CSP:
 - structured and simple representation
 - general purpose algorithms

Constraint Satisfaction Problem

- Defined by
 - n variables X_i which define a state
 - Each variable has a domain D_i of possible values
 - m constraints C_j
 - Each constraint involves some subset of variables
 - Specifies the allowable combinations of values
- A state of the problem: assignment of values to some or all of X_i s

Constraint Satisfaction Problem

- **Consistent** or legal assignment does not violate constraints
- **Complete assignment** that satisfies all constraints is a solution
- A complementary objective function may be defined

Example: Map-Coloring



Variables WA, NT, Q, NSW, V, SA, T

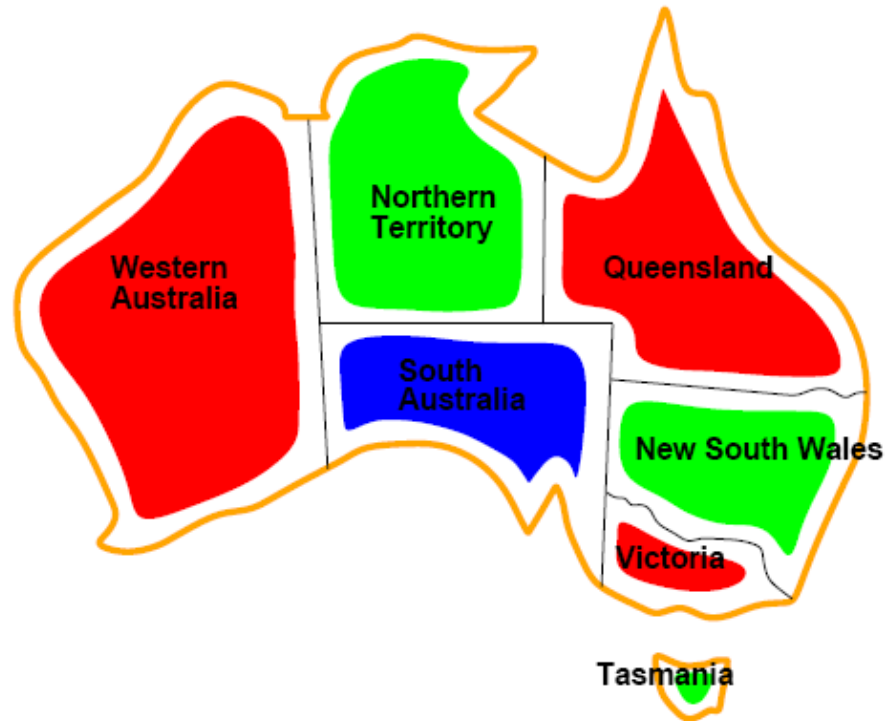
Domains $D_i = \{red, green, blue\}$

Constraints: adjacent regions must have different colors

e.g., $WA \neq NT$ (if the language allows this), or

$(WA, NT) \in \{(red, green), (red, blue), (green, red), (green, blue), \dots\}$

Example: Map-Coloring




Solutions are assignments satisfying all constraints, e.g.,

$\{WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green\}$

Constraint graph

- Constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the graph structure
 - to speed up search. e.g., Tasmania



Standard search formulation

- States are defined by the values assigned so far
 - Initial state: the empty assignment, $\{\}$
 - Successor function: assign a value to an unassigned variable that does not conflict with the current assignment
 - fail if no legal assignments (not fixable!)
 - Goal test: the current assignment is complete
 - Path cost: a constant cost for every step

Standard search formulation

- The search formulation is the same for all CSPs!
- What is the depth of solution?
 - Which type of search?
- Path is irrelevant, so can also use complete-state formulation
- The number of leaves! vs possible assignments
 - Commutativity
 - Consider only a single variable at each node

Varieties of CSPs

- Discrete variables
 - finite domains; size $d \rightarrow O(d^n)$ complete assignments
 - e.g., Boolean CSPs including Boolean satisfiability (NP-complete)
 - infinite domains (integers, strings, etc.)
 - e.g., job scheduling, variables are start/end days for each job
 - need a constraint language, e.g., $StartJob1 + 5 \leq StartJob3$
 - linear constraints solvable, nonlinear undecidable

Varieties of CSPs

- Continuous variables
 - e.g., start/end times for Hubble Telescope observations
 - linear constraints solvable in polynomial time by LP methods

Varieties of constraints


- Unary constraints involve a single variable
 - $SA \neq \text{green}$
- Binary constraints involve pairs of variables
 - $SA \neq WA$
- A binary CSP has only binary constraints, constraint graphs

Varieties of constraints

- Global constraints involve an arbitrary number of variables: cryptarithmic column constraints
 - constraint hypergraph
 - can be reduced to binary constraints
- Preferences (soft constraints)
 - red is better than green
 - often encoded using costs against the overall objective function
 - constrained optimization problems

Example: Cryptarithmic

$$\begin{array}{r} \text{T W O} \\ + \text{T W O} \\ \hline \text{F O U R} \end{array}$$



Variables: $F T U W R O X_1 X_2 X_3$

Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

Constraints

$alldiff(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1$, etc.

Real-world CSPs

- Assignment problems
 - who teaches what class
- Timetabling problems
 - which class is offered when and where?
- Hardware configuration
- Spreadsheets
- Transportation scheduling
- Factory scheduling
- Floorplanning

Benefits of modeling as a CSP

- Representation of states conforms to a standard pattern
- The successor function and goal test can be written in a generic way
- Devising generic heuristics
- The structure of the constraint graph can be used to simplify the solution process

Backtracking search

- Variable assignments are commutative
 - [WA=red then NT =green] same as [NT =green then WA=red]
- Only need to consider assignments to a single variable at each node
 - $b=d$ and there are d^n leaves

Backtracking search

- Depth-first search with single-variable assignments
- The basic uninformed algorithm for CSPs
- Can solve n-queens for $n \approx 25$

Backtracking search


```
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return BACKTRACK({ }, csp)

function BACKTRACK(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
  var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment then
      add {var = value} to assignment
      inferences  $\leftarrow$  INFERENCE(csp, var, value)
      if inferences  $\neq$  failure then
        add inferences to assignment
        result  $\leftarrow$  BACKTRACK(assignment, csp)
        if result  $\neq$  failure then
          return result
      remove {var = value} and inferences from assignment
  return failure
```


Backtracking example




Backtracking example



Backtracking example



Backtracking example




Improving backtracking efficiency

- General-purpose methods can give huge gains in speed:
 - Which variable should be assigned next?
 - In what order should its values be tried?
 - Can we detect inevitable failure early?
 - Can we take advantage of problem structure?

Minimum remaining values


- Minimum remaining values (MRV):
 - choose the variable with the fewest legal values



- Most constrained variable, fail-first heuristic



Degree heuristic

- Tie-breaker among MRV variables
- Degree heuristic:
 - choose the variable with the most constraints on remaining variables



Least constraining value


- Given a variable, choose the least constraining value:
 - the one that rules out the fewest values in the remaining variables



- Combining these heuristics makes 1000 queens feasible


Forward checking

- Idea: Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



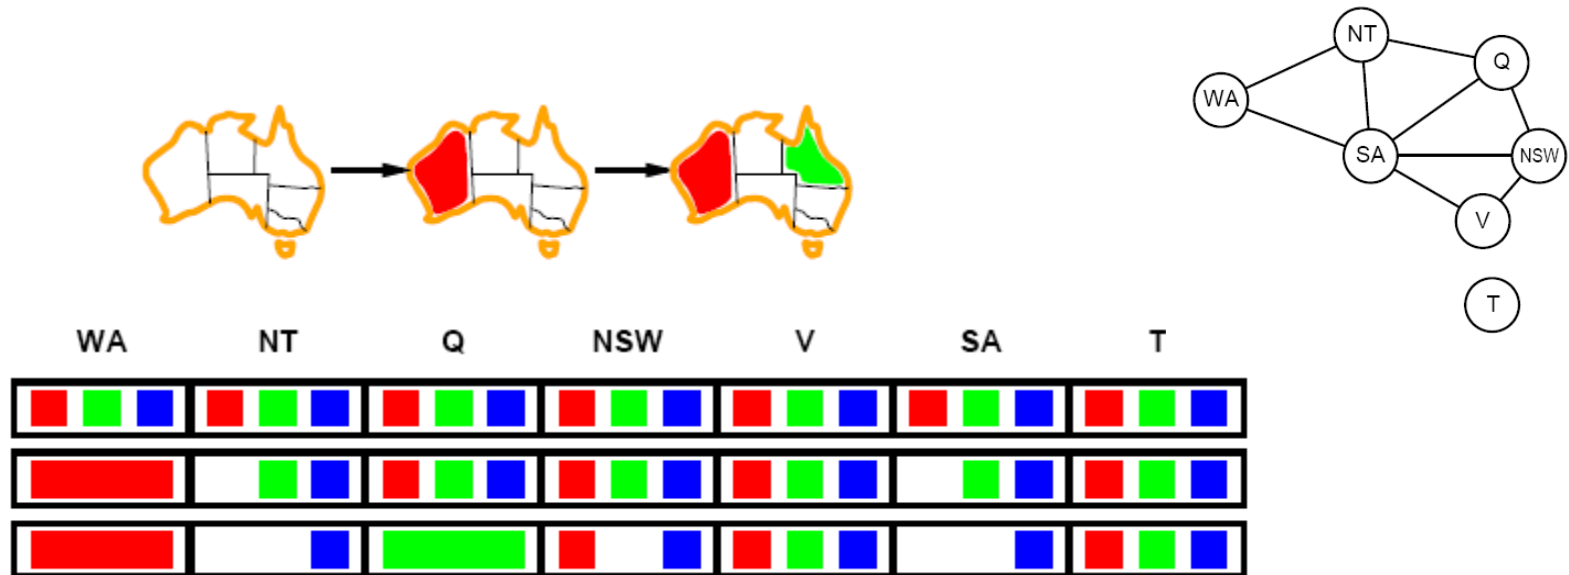
Forward checking

- Idea: Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



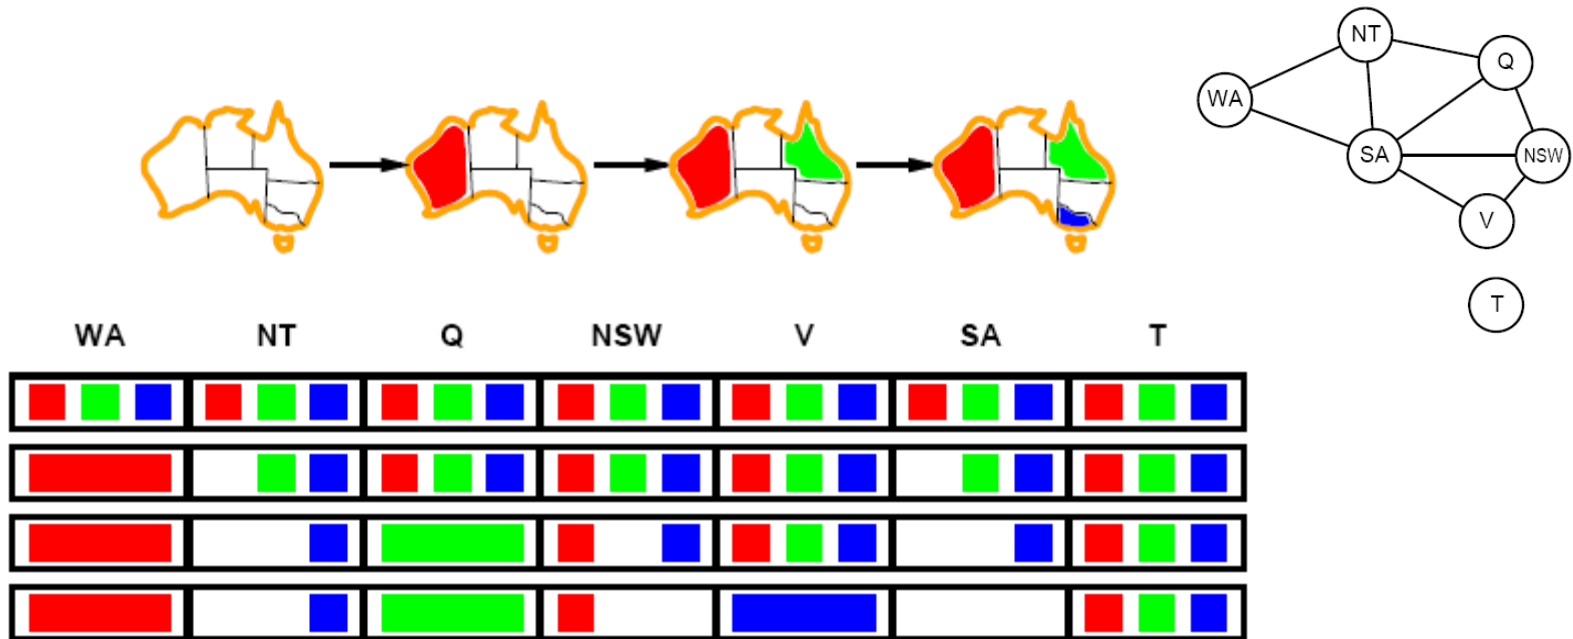
Forward checking

- Idea: Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values





Forward checking

- Idea: Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



Constraint propagation



- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures



Arc consistency

- Simplest form of propagation makes each arc consistent



$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



Arc consistency

- Simplest form of propagation makes each arc consistent


$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y




Arc consistency

- Simplest form of propagation makes each arc consistent

$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y




If X loses a value, neighbors of X need to be rechecked




Arc consistency

- Simplest form of propagation makes each arc consistent

$X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



If X loses a value, neighbors of X need to be rechecked



Arc consistency algorithm

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X , D , C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

 (X_i , X_j) \leftarrow REMOVE-FIRST(*queue*)

if REVISE(*csp*, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** X_i .NEIGHBORS - $\{X_j\}$ **do**

 add (X_k , X_i) to *queue*

return true

function REVISE(*csp*, X_i , X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow false

for each x **in** D_i **do**

if no value y in D_j allows (x,y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i


revised \leftarrow true

return *revised*

k-consistency

- Arc-consistency does not reveal every possible inconsistency
- k-consistency:
 - for any $k-1$ variables and
 - for any consistent assignment
 - A consistent value can be assigned to any k^{th} variable
- 1-consistency: node consistency
- 2-consistency: arc consistency
- 3-consistency: path consistency
- A graph is strongly k-consistent if it is k-consistent and also $(k-1)$ -consistent, ... 1-consistent

Problem structure



Problem structure

Suppose each subproblem has c variables out of n total

Worst-case solution cost is $n/c \cdot d^c$, **linear** in n


E.g., $n = 80$, $d = 2$, $c = 20$

$2^{80} = 4$ billion years at 10 million nodes/sec

$4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec


Tree-structured CSPs

- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$



Algorithm for tree-structured CSPs

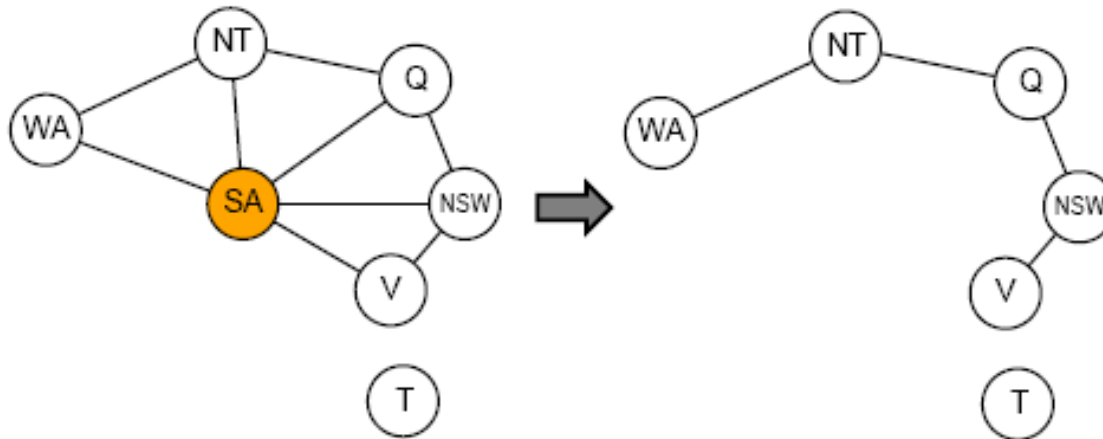
1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering



2. For j from n down to 2 , apply $\text{REMOVEINCONSISTENT}(\text{Parent}(X_j), X_j)$
3. For j from 1 to n , assign X_j consistently with $\text{Parent}(X_j)$

Nearly tree-structured CSPs

Conditioning: instantiate a variable, prune its neighbors' domains



Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

Cutset size $c \Rightarrow$ runtime $O(d^c \cdot (n - c)d^2)$, very fast for small c